



CS 5/7320
Artificial Intelligence

Search with Uncertainty

AIMA Chapters 4.3-4.5

Slides by Michael Hahsler
with figures from the AIMA textbook





This work is licensed under a [Creative Commons Attribution-ShareAlike 4.0 International License](https://creativecommons.org/licenses/by-sa/4.0/).



Online Material


Recap: Solving Search Problems under Certainty



No Uncertainty

- **Full observability:** The agent always knows (=can observe) the state.
- **Deterministic environment** with a known transition model
 $Result(s, a) = s'$
The agent can predict the outcome of its actions.

State space: A state completely describes the condition of the environment and the agent.



Solution: Use tree search in the planning phase to create a **sequence of actions** also called a **plan**. Then blindly execute the plan: **[Suck, Right, Suck]**

Sources and Consequence of Uncertainty

Sources: The environment may be

- **Not fully observable:** The agent may be uncertain about its current state.
- **Stochastic (transition function):** The agent may not be able to perfectly predict the outcome of its actions.

Consequences:

1. The agent needs to keep track of all the states it could be in. This set is called a ***belief state***.
2. A fixed precomputed plan (sequence of actions) does not work for stochastic transition functions, but a ***conditional plan (also called strategy or policy)*** that depends on percepts is needed.

Types of uncertainty in the environment*



Nondeterministic Actions:

Outcome of an action in a state is uncertain.



No observations:

Sensorless problems.



Partially observable environments:

The agent cannot directly observe the state of the environment.



Exploration:

Unknown environments and online search.

* we will quantify uncertainty with probabilities later.


A dynamic background image showing a bright yellow powder or smoke explosion against a black background. The particles are concentrated on the right side and spread out towards the left, creating a sense of motion and energy.

Nondeterministic Actions

Stochastic Environment (Stochastic Transition Model)

Definition: Nondeterministic Actions

The outcome of actions in the environment is nondeterministic = the **transition model needs to describe uncertainty.**




Note the 's' here

Example transition:

$$Results(s_1, a) = \{s_2, s_4, s_5\}$$

i.e., action a in s_1 can lead to one of several states.


Example: Erratic Vacuum World



Regular fully observable, deterministic vacuum world, but the action '**suck**' is more powerful and **nondeterministic**:


- a) **On a dirty square:** cleans the square and sometimes cleans dirt on adjacent squares as well.
- b) **On a clean square:** sometimes deposits some dirt on the square.

Example: Erratic Vacuum World




Suck can lead to two different states! We need a conditional plan
[Suck, **if** State = 5 **then** [Right, Suck] **else** []]

Transition Model as an AND-OR Search Tree



LOOP: No need to continue search.
Solution is the same as above.

Search the AND-OR Tree



- **Goal:** Find a subtree with one action for each OR node and considering all outcomes of the AND nodes that has only goal leaf nodes.
- Descend the tree depth-first:
 - OR node: trying one action at a time.
 - AND node: consider all outcomes and check recursively.
 - Ignore cycles.
 - Abandon a subtree if not all leaf nodes are the desired goal nodes.
 - Stop when **the first complete subtree with only goal leaf nodes is found.**
- Construct the conditional plan that represents the subtree starting at the root node.

Conditional Plan:
[Suck, if State = 5 then [Right, Suck] else []]

Recursive AND-OR Tree Search (DFS)

= nested If-then-else statements

```
function AND-OR-SEARCH(problem) returns a conditional plan, or failure  
  return OR-SEARCH(problem, problem.INITIAL, [])
```

path is used for cycle checking!

```
function OR-SEARCH(problem, state, path) returns a conditional plan, or failure  
  if problem.IS-GOAL(state) then return the empty plan  
  if IS-CYCLE(path) then return failure // don't follow loops using path  
  for each action in problem.ACTIONS(state) do // try all possible actions  
    plan  $\leftarrow$  AND-SEARCH(problem, RESULTS(state, action), [state] + path)  
    if plan  $\neq$  failure then return [action] + plan  
  return failure // fail means we found no action that leads to  
                // a goal-only subtree  
  
function AND-SEARCH(problem, states, path) returns a conditional plan, or failure  
  for each si in states do // consider all possible outcomes, none can fail!  
    plani  $\leftarrow$  OR-SEARCH(problem, si, path)  
    if plani = failure then return failure // fail if we find any non-goal subtree  
  return [if s1 then plan1 else if s2 then plan2 else ... if sn-1 then plann-1 else plann]
```


Notes:

- The DFS search tree is implicitly created using the call stack (recursive algorithm).
- DFS is **not optimal**! It returns the first valid plan (subtree) that it finds.

An Agent using the Conditional Plan

- Planning uses search to find a conditional plan that always leads to a goal state.
- The conditional plan can be executed by a **model-based reflex agent** that uses a program counter to execute the plan and **percepts** for the conditions in the if-statements.

Example: After the initial action “suck”



A futuristic robot with a metallic, grey and black body is the central figure. It wears a black blindfold over its eyes. The robot's head is encased in a helmet-like structure with some wires or sensors extending from the sides. The background is a dimly lit, industrial or laboratory-like environment with various cables and a small screen visible on the right side. The overall tone is dark and mysterious.

Search With No Observations

Using actions to “coerce” the world into a smaller set of known states

Sensorless Problems


Conformant problem: The agent has no sensors, so the environment is not observable.

Why is this useful?



- **Example:** Doctor prescribes a broad-band antibiotic instead of performing time-consuming blood work to find a more targeted antibiotic. This saves time and money.
- **Basic idea:** Find a solution (a plan) that **works (reasonably well) from any state** and then just blindly execute it.




Definition: Belief State



- The agent does not know exactly what state it is in.
- However, it may know that it is in one of a set of possible states. This set is called a **belief state** of the agent.
- Example: $b = \{s_2, s_4, s_6\}$




8






Actions to Coerce the World into Known States




- Actions can reduce the number of possible states.
- **Example:** Deterministic but unobservable vacuum world. The agent does not know its position or the dirt distribution.

Initial belief state {1,2,3,4,5,6,7,8}






Actions to Coerce the World into Known States




- Actions can reduce the number of possible states.
- **Example:** Deterministic but unobservable vacuum world. The agent does not know its position or the dirt distribution.






Actions to Coerce the World into Known States





- **Plan:** The action sequence [right, suck, left, suck] coerces the world into the goal state 7. This plan works from any initial state!
- **Note:** There are no observations, so there is no need for a conditional plan.





The Reachable Belief State Space



The size of the belief state space is the powerset of the original N states:

$$\mathcal{P}_s = 2^N = 2^8 = 256$$

Only a small fraction (12 belief states) are reachable by actions.

No observations, so we get a solution sequence from an initial belief state:
[Right, Suck, Left, Suck]

Finding a Plan

Note: State space size makes this impractical for larger problems!

Formulate as a regular search problem and solve with DFS, IDS, BFS or A*:

- **States:** All belief states (=powerset $\mathcal{P}_s$ of the set of N states has size 2^N)
- **Initial state:** Often the belief state containing all states.
- **Actions:** Available actions of a belief state are the union of the possible actions for all the states it contains.
- **Transition model:** $b' = Results(b, a) = \{s' : s' = Result(s, a) \text{ and } s \in b\}$
- **Goal test:** Does the belief state only contain goal states?
- **Simplifying property:** If a belief state (e.g., $b_1 = \{1,2,3,4,5\}$) is solvable (i.e., there is a sequence of actions that coerce all states to only goal states), then belief states that are subsets (e.g., $b_2 = \{2,5\}$) are also solved using the same action sequence. This can be used to prune the search tree.

Other approach:

- **Incremental belief-state search.** Generate a solution that works for one state and check if it also works for all other states. If it does not, then modify the solution slightly. This is similar to local search.


Case Study

The agent can move up, down right, and left.
The agent has **no sensors** and does not know its current location.

1. Can you navigate to the goal location?
How?

2. What would you need to know about the environment?

3. What type of agent can do this?





Partially Observable Environments

Using Observations to Learn About the State


Percepts and Observability

- Many problems cannot be solved efficiently without sensing (e.g., 8-puzzle).
- We need to see at least one square.

Percept function: $Percept(s)$

... s is the state


- **Fully observable:** $Percept(s) = s$
- **Sensorless:** $Percept(s) = None$
- **Partially observable:** $Percept(s) = o$
 o is called an observation and tells us something about s



$Percept(s) = Tile7$

Problem: Many states (different order of the hidden tiles) can produce the same observation!

Use Observations to Learn About the State



The agents chooses an action and then receive an observation.

Idea: Observations can be used to learn about the agent's state.

Assume we have a current belief state b (i.e., the set of states we could be in).

1. Prediction for action: Choose an action a and compute a new belief state that results from the action using the transition model.

$$\hat{b} = \text{Predict}(b, a) = \bigcup_{s \in b} \text{Result}(s, a)$$


2. Update with observation: You receive an observation o and only keep states that are consistent with the new observation. The filtered belief after observing o is:

$$b_o = \text{Update}(\hat{b}, o) = \{s : s \in \hat{b} \wedge \text{Percept}(s) = o\}$$



Writing both steps as one update:

$$b \leftarrow \text{Update}(\text{Predict}(b, a), o)$$

Example: Deterministic local sensing vacuum world




Only local dirt sensor



$$b \leftarrow \text{Update}(\text{Predict}(b, a), o)$$



$$\text{Update}(\text{Predict}(\{1,3\}, \text{Right}), [B, \text{Dirty}]) = \{2\}$$

Solving Partially Observable Problems




Only local dirt sensor

Use an AND-OR tree of belief states to create a **conditional plan**.




Solving Partially Observable Problems 2




Only local dirt sensor

Use an AND-OR tree of belief states to create a **conditional plan**.




Plan: [**Suck**, ...]

Use an AND-OR tree of belief states to create a **conditional plan**.




Solving Partially Observable Problems 4




Only local dirt sensor

Use an AND-OR tree of belief states to create a **conditional plan**.



$b = \{6\}$ is the result of the update with $o = [B, Dirty]$

State Estimation and Approximate Belief States




- Agents choose an **action** and then receive an **observation** from the environment.
- The agent keeps track of its belief state using the following update:

$$b \leftarrow \text{Update}(\text{Predict}(b, a), o)$$

- This process is often called
 - **monitoring**,
 - **filtering**, or
 - **state estimation**.
- **Issue:** The agent needs to be able to update its belief state following observations in **real time**! For many practical applications, there is only time to compute an **approximate belief state**! Such approximations are commonly used in control theory and reinforcement learning.

Environment State vs. Agent State vs. Model


- The complete factored **environment state** description represents the true state of the world and is necessary and sufficient to apply the transition function. E.g.,



- **Full observability:** Agent does not need to store the environment state since it is the percept. The agent may **store additional information** it uses to decide on the next action. E.g.,
 - The step in its internal “program” used to decide on the next action.
 - Cleaning robot: Cleaning mode vs. navigating to the charging station.
- **Partial observability:** Agent has an **agent state** representing what the agent knows and uses it to decide on actions. The agent uses actions and observations to update its agent state to estimate the environment state. E.g.,
 - Remember prior observations that are not in the current percept: Tile A was clean at the last visit.
 - Estimates for unobserved fluents.
 - The agent’s current belief state.
- **Model:** Is the agent’s internal knowledge about how the environment behaves. It can range from:
 - A simple internal variable: Cleaning mode vs. navigating to the charging station. This implicitly models the environment’s behavior that the battery will run out.
 - A full transition model that can predict the next state given a state and a possible action. This can be used for planning.
 - Information used to estimate the current state from partial observations.

Case Study:

Partially
Observable
8-Puzzle




Partially Observable 8-Puzzle


Give a problem description for this problem.

- States:
- Initial state:
- Actions:
- Transition model:
- Goal test:
- Percept function:

This problem can be solved using an AND-OR Tree, but is there an easier solution?

- a. What type of agents would we use?
- b. What algorithms can be used?





Exploration

Unknown Environments and Online Search



Recap: Offline Search

- **Offline search aka planning:** Create a plan using the state space and the transition model before taking any action.
- The **plan** can be
 - a **sequence of actions**, or
 - a **conditional plan** that uses observations to account for uncertainty or imperfect observability.
- The agent plans using search with the known transition function to predict the consequence of actions.
- **Issue:** In an **unknown environment**, we do not know the transition function.
- We cannot predict outcomes of actions; therefore, we cannot plan using offline search!

Online Search

- **Online search** does not use planning! It explores the real world one action at a time. Offline prediction and update are replaced by “act” and “observe.”



- Useful for
 - **Unknown environment:** The agent has no complete model of how the environment works. It needs to explore an unknown state space and/or what actions do. I.e., it needs to **learn the transition function**
$$f : S \times A \rightarrow S$$
 - **Real-time problems:** When offline computation takes too long, and there is a penalty for sitting around and thinking.
 - **Nondeterministic domain:** Conditional plans become very large. Only focus on what happens instead of planning for everything!

Design Considerations for Online Search

- **Knowledge:** What does the agent already know about the outcome of actions? E.g.,

- Does go north and then south lead to the same location?
- Where are the walls in the maze?



Transition
function

Often a part or all of the transition function is unknown!

- **We need a safely explorable state space/world:** There are **no irreversible actions** (e.g., traps, cliffs) or the agent needs to be able to avoid these actions during exploration using percepts.
- **Exploration order is important:** Expanding nodes in **local order** (= close by) is more efficient if you must execute the actions to get observations: Use depth-first search with backtracking instead of BFS or A* Search.

Online Search: A Model-based Reflex Agent to Learn the Transition Model

Setting: Environment is deterministic and fully observable (= the percept is the full state) but the transition model function *result()* is unknown.

Approach: The agent builds the map $result(s, a) \rightarrow s'$ by trying all actions and backtracks when all actions in a state have been explored (this is a form of iterative DFS called backtracking DFS).

function ONLINE-DFS-AGENT(*problem*, *s*) **returns** an action

s, *a*, the previous state and action, initially null

persistent: *result*, a table mapping (*s*, *a*) to *s'*, initially empty

untried, a table mapping *s* to a list of untried actions

unbacktracked, a stack with the current path

if *problem*.IS-GOAL(*s'*) **then return** *stop*

if *s'* is a new state (not in *untried*) **then** *untried*[*s'*] $\leftarrow$ *problem*.ACTIONS(*s'*)

if *s* is not null **then**

result[*s*, *a*] $\leftarrow$ *s'*

add *s* to the front of *unbacktracked*[*s'*]

if *untried*[*s'*] is empty **then**

if *unbacktracked*[*s'*] is empty **then return** *stop*

else *a* $\leftarrow$ an action *b* such that *result*[*s'*, *b*] = POP(*unbacktracked*[*s'*])

else *a* $\leftarrow$ POP(*untried*[*s'*])

s $\leftarrow$ *s'*

return *a*

Learn the result function
(= transition function)

Untried is the "frontier"

Unbacktracked stores the current path


Record found transitions

Keep breadcrumbs to go back
later

Use breadcrumbs to walk back

Case Study: DFS with Backtracking for an unknown Maze

- We don't have a map (transition function) of the maze. We can only see adjacent squares.
- We cannot plan, so we must explore by walking around!
- A simple method is **backtracking DFS** that only stores the current path for backtracking (on a stack) to get back to untied actions when we run into a dead end (think leaving breadcrumbs or a string).
- This is an iterative implementation of DFS without a reached data structure. Unbacktaced represents the currently explored path, and untried represents the frontier. DFS memory management applies.





Important concepts that you should be able to explain and use now...

- Difference between solution types:
 - a. a fixed action sequence (a plan),
 - b. a **conditional plan** (also called a strategy or policy), and
 - c. **exploration**.
- What are **belief states**?
- How actions can be used to coerce the world into known states.
- How actions and observations (from **percept functions**) can be used to learn about the state: State estimation with repeated predict and update steps.
- The use of AND-OR trees to solve small problems.
- Large problems are hard!